


Forecasting Brazil's Next Tailings-Dam Disaster


*Sentinela — a calibrated hierarchical-Bayesian forecaster
for Brazilian mine-tailings dam failures*

Sentinela — predicted 12-month tailings-dam failure-risk field over Brazil
spike height scales with failure probability · teal-violet-magenta by risk decile · 877 active dams

(A) National failure-risk field



(B) Minas Gerais iron-ore corridor



Predicted 12-month failure-risk field over Brazil's terrain. (A) national view; (B) Minas Gerais iron-ore corridor with the reference-failure complexes labelled. Spike height scales with failure probability; colour runs teal → violet → magenta by risk decile.

Viktor Smirnov

Arakon

24 May 2026 • pre-print

Open data • open code • reproducible from a clean clone
github.com/ViktorSmirnov71/sentinela-bayes-br

Contents

Abstract	1
1. Introduction	1
2. Related work	2
2.1 Single-event precursor reconstruction	2
2.2 Generalised early-warning and ML on InSAR	2
2.3 Population-scale risk classification	2
2.4 The gap	2
3. Data	2
3.1 Cohort: SIGBM	3
3.2 Labels: Brazilian failure events	3
3.3 Precursor signal: Sentinel-1 InSAR	3
3.4 Forcing variables: rainfall and antecedent moisture	4
3.5 Exposure: IBGE	4
4. Methodology	4
4.1 Formal task	4
4.2 Why a hierarchical-Bayesian model, not a discriminative learner	4
4.3 Validation protocol	5
4.4 Reproducibility	5
5. Results	5
5.1 Construction-method posterior	6
5.2 Within-upstream ranking	6
5.3 Calibration and discrimination	7
5.4 Fundão retrospective: a partial-precursor finding	7
5.5 Brumadinho B1 retrospective: a partial-precursor hit	8
5.6 Visualisation	9
6. Discussion	9
6.1 What we learn	9
6.2 What we do not learn	10
6.3 Ethics	10
7. Limitations and threats to validity	10
8. Future work	10
9. Conclusion	11
Figures and tables	11
Acknowledgements	11
References	12

Appendix A — Figures

13

Appendix B — Tables

16

Abstract

Catastrophic failure of mine-tailings storage facilities (TSFs) in Brazil killed at least 289 people in two events in the last decade (Mariana 2015, Brumadinho 2019). After Brumadinho, three things became reproducibly true: the Brazilian regulator publishes a national TSF registry (SIGBM, 911 dams in our 2026 snapshot); the European Copernicus programme freely distributes Sentinel-1 SAR with a 6–12 day revisit over all of Brazil since 2014; and post-mortem analyses (Grebby et al. 2021; Carlà et al. 2019) confirmed that the B1 Córrego do Feijão collapse was detectable from Sentinel-1 InSAR at least five months in advance. Yet there is no published, openly reproducible, per-dam probabilistic forecaster that uses these inputs.

We present **Sentinela**, a hierarchical-Bayesian model that produces a calibrated 12-month-ahead failure probability for every active Brazilian mine-tailings dam from purely public data. The model composes a literature-informed Beta-Binomial construction-method posterior, a James–Stein-shrunk per-operator random effect, bounded logit shifts for engineering features (height, volume, age, regulator risk category), and a bounded InSAR-precursor term derived from Sentinel-1 line-of-sight displacement. We report three results. (1) On the 2026 SIGBM snapshot the engineering-feature ranking independently surfaces the dams ANM has separately placed under emergency declarations: the single Emergency-Level-3 dam in the country (Barragem Serra Azul) and the entire Level-2 Forquilha cluster appear in the top-15, despite the model receiving no emergency-level input — an external validation that the within-construction-class ordering is meaningful. (2) We process 148 Sentinel-1 interferograms through ASF HyP3 at the actual coordinates of both reference failures and run a month-by-month retrospective; at Brumadinho B1 the predicted risk spikes in April 2018, inside the “definitive but emergent risk” window reported independently by Grebby et al. (2021), a partial replication of the published precursor with a purely open pipeline. (3) We document a “registry-amnesia” bias (failed dams are struck from SIGBM post-failure) and a cross-relative-orbit SBAS pairing failure mode (100 % of our HyP3 failures), both of which are methodologically instructive at population scale. We argue that in the one-linked-positive regime that characterises this problem, a transparent literature-prior Bayesian model is the honest choice over discriminative learners, which we show empirically either memorise the single positive or collapse to zero. All code, data, figures, the 3D visualisation, and this manuscript are openly released and reproducible from a clean clone in five commands.

1. Introduction

A tailings dam is the engineered wall that retains the slurry by-product of mineral extraction. Catastrophic failure releases the impounded volume as a flowing mudwave. Mariana 2015 destroyed Bento Rodrigues and contaminated 600 km of the Doce river. Brumadinho 2019 killed approximately 270 people in 90 seconds. These are not isolated events: the World Mine Tailings Failures catalogue records eight Brazilian severity-4-or-higher events since 1986.

Three developments after Brumadinho make per-dam quantitative forecasting newly tractable:

1. **SIGBM** (Sistema Integrado de Gestão de Barragens de Mineração) — the Agência Nacional de Mineração’s national TSF registry, public since 2019 [[@anm_sigbm](#)].
2. **Sentinel-1** — ESA Copernicus C-band SAR, full archive open since 2014, 6–12-day revisit, millimetre-precision interferometric ground-deformation measurement.

3. **Post-mortem replication studies** [grebby2021; carla2019; du2020; gama2023] that confirm the B1 collapse left a clear InSAR signature months before failure.

The literature has progressed from “*InSAR can post-hoc identify precursors at known failures*” to “*InSAR + ML can flag elevated risk at active facilities*”. No published model jointly (a) uses SIGBM as the cohort, (b) propagates calibrated uncertainty, and (c) is reproducibly open from data to forecast. This paper fills that gap with the simplest methodology that does so.

2. Related work

2.1 Single-event precursor reconstruction

grebby2021 applied advanced InSAR analysis to Sentinel-1 over B1 and detected ground-deformation precursors ≥ 5 months before the 25 January 2019 collapse. carla2019 independently derived an inverse-velocity time-of-failure prediction. du2020 performed an InSAR time-series risk assessment of the Brumadinho cluster. gama2023 published a slip-surface mechanism account attributing the failure to delayed static liquefaction. mura2018 used DInSAR with TerraSAR-X to monitor the surviving dikes at the Germano complex post-Fundão. These works establish that the geophysical precursor signal exists in publicly-available Sentinel-1 data at known failures.

2.2 Generalised early-warning and ML on InSAR

macchiarulo2023 reduced false alarms by combining InSAR deformation with SAR-derived moisture content. manconi2024 catalogued practical considerations for operational Sentinel-1 monitoring of tailings dams. zhang2025 proposed an InSAR framework that robustly separates consolidation settlement from failure-relevant deformation. zheng2025 integrated InSAR with 3D numerical models. None of these output calibrated per-dam probabilities at population scale, none jointly use construction-type metadata, and none are reproducibly open from data to forecast.

2.3 Population-scale risk classification

delima2022 classified Brazilian TSFs by risk using SIGBM-derived features (without InSAR), but treats ANM’s administrative risk score as the target, not failure events. rana2021 published a global statistical model of tailings-dam failure frequency vs. construction type using WMTF data — a population baseline rate, not a per-dam forecast.

2.4 The gap

Combining the three lines: - Post-hoc precursor work shows the *signal exists*. - Methodological InSAR-ML work shows the *signal can be extracted*. - Population-scale work establishes a *cohort* but uses administrative scores as labels.

What does not exist: a reproducible per-dam probabilistic forecaster for Brazilian TSFs that (i) uses SIGBM as its cohort, (ii) ingests Sentinel-1 InSAR-derived precursor features, (iii) is trained against documented failure or near-miss events with explicit survival censoring, and (iv) is evaluated with proper uncertainty quantification. This is the target of Sentinela.

3. Data

We assemble the analysis from five categories of public open data; full provenance, schema, and access mechanisms are documented in docs/03-data.md and data/README.md of the open-source repository.

3.1 Cohort: SIGBM

The Agência Nacional de Mineração’s public web app at app.anm.gov.br/sigbm/publico provides a manual CSV export with 911 dams (2026-05-21 snapshot). Schema includes construction method (upstream / downstream / centerline / single-stage / dyke / unknown), current height and volume, primary ore, operator CNPJ, the ANM Risk Category (CRI) and Damage-Potential Category (DPA), declared emergency level, and operational status. The export is semicolon-delimited UTF-8-with-BOM with Brazilian decimal format and DMS coordinates; we canonicalise these in `src/sentinela/io/sigbm.py`.

Snapshot composition. The cohort spans 21 states, dominated by Minas Gerais (35%), Mato Grosso (20%), and Pará (13%). Construction methods are distributed 53% single-stage, 24% downstream, 13% centerline, 5% upstream (the historically dangerous class with 45 dams), 5% unknown. As of the snapshot, 822 dams are at no declared emergency level, 12 at Alert level, 69 at Emergency Level 1, 7 at Emergency Level 2, and 1 at Emergency Level 3 — the single most-immediately-concerning structure currently in Brazil.

Methodology note — registry amnesia. ANM removes failed dams from the active SIGBM registry post-failure. Vale’s B1 Córrego do Feijão (2019), Samarco’s Fundão (2015), and Herculano’s B1 (2014) are all absent from the current export. The Fundão event has a credible proxy via `dam_id` 8765, *Barragem de Germano* — the surviving Samarco structure at the same mining complex, now in decommissioning. Other historical failures need pre-failure SIGBM snapshots (obtainable from Wayback Machine archives or by FOI request) to be modelled prospectively. This biases the cohort toward “dams that have not yet failed” and is the single biggest limitation of the present work.

3.2 Labels: Brazilian failure events

A hand-curated, citation-anchored table of eight documented Brazilian TSF-failure events from 1986 to 2022 is committed to the repository at `data/external/brazilian_failures.csv`. The table preserves the Bowker–Chambers severity coding [`@wmtf`]; the primary cutoff for the positive class is $\text{severity} \geq 4$.

3.3 Precursor signal: Sentinel-1 InSAR

Sentinel-1 SLC scenes intersecting each dam are searched via the ASF API (`asf_search`). For each dam we generate the short-baseline-subset pair list, submit InSAR processing jobs to ASF’s HyP3 hosted service [`hyp3_sdk`], and download the unwrapped-phase GeoTIFFs as they complete. The HyP3 path avoids 10+ GB local SLC downloads and removes the need for a working ISCE2 / GAMMA toolchain. Across the project we processed three HyP3 batches totalling 276 submitted jobs (148 succeeded): the Germano proxy (`dam_id` 8765, 42 pairs), the actual Fundão coordinates (42 pairs), and the actual Brumadinho B1 coordinates (192 submitted, 64 succeeded — see §5.5 on the cross-track failure mode). Each succeeded job yields one unwrapped-phase raster which we sample at the dam centroid.

The orbit auto-selection step is non-trivial. The first manual submission used the script’s `ASCENDING` default and returned zero scenes; we verified that early-mission Sentinel-1 coverage of Minas Gerais is descending-only. We now default to a per-dam `pick_best_orbit` lookup that searches both orbits and uses whichever has more SLCs in the target window.

3.4 Forcing variables: rainfall and antecedent moisture

Three sources of rainfall data will be ingested in the next experimental phase: CHIRPS daily 0.05° rainfall (1981–present), INMET BDMEP station records (1961–present), and ERA5-Land hourly reanalysis. We compute SPI-3 and SPI-6 z-scores against the local long-term climatology — raw rainfall is misleading across Brazil’s climatic variation, and the standardised indices are the established operationalisation of “wet” / “dry” relative to a location’s baseline.

3.5 Exposure: IBGE

For translating failure probability into expected affected population, we will rasterise IBGE Setor Censitário population data onto a common grid and compute the population within a runout-distance buffer for each dam. This data does not enter the forecasting model directly; it informs decision-curve analysis when the model is evaluated as an operational input.

4. Methodology

4.1 Formal task

Let $\mathcal{D} = \{d_1, \dots, d_N\}$ index the active dams in SIGBM. For each (dam, month) pair (d_i, t) we model

$$P(F_{i,[t,t+H]} = 1 \mid x_{i,\leq t})$$

where $F_{i,[t,t+H]}$ is the binary event “dam d_i experiences a failure of severity ≥ 4 within horizon H from time t ”, with $H = 12$ months as the primary target horizon. The feature vector $x_{i,\leq t}$ comprises static engineering metadata, InSAR-derived precursor features computed over the trailing 12 months, antecedent climate variables, and discrete operational-status indicators.

4.2 Why a hierarchical-Bayesian model, not a discriminative learner

With one linked positive event in the current cohort (Fundão, mapped via proxy to dam_id 8765), discriminative models either memorise the single positive’s unique feature signature or smooth to zero everywhere. We documented both failure modes empirically:

- **LightGBM with min_data_in_leaf=20:** training AUROC 0.9995, ECE 0.001, but assigned all non-target dams $\approx 0\%$ predicted risk. Useless.
- **LightGBM with min_data_in_leaf=500:** all probabilities ≈ 0 , AUROC collapsed to 0.18. Also useless.

The honest model in the small-positive regime is a **hierarchical Bayesian estimator** with informative priors. We compose:

1. **Level 1** — Beta-Binomial smoothed per-construction-method failure rate. For each class c :

$$\hat{p}_c = \frac{k_c + \alpha p_c^{\text{prior}}}{n_c + \alpha}$$

where p_c^{prior} comes from @rana2021 and the Bowker–Chambers catalogue (annual rates: up-stream 0.50%, centerline / single-stage / dyke / unknown 0.10%, downstream 0.05%) and $\alpha = 10,000$ equivalent dam-months.

2. **Level 2** — per-operator logit shift via James–Stein-style shrinkage, capped at ± 1 logit unit so a single positive cannot make any operator look like an outlier.
3. **Level 3** — bounded logit shifts for continuous engineering features: height, volume, age (z-scored against the cohort), and the regulator-assigned CRI ordinal. Coefficients are conservative (0.10–0.30 in logit units) and informed by the published effect sizes in @rana2021.

The Level-1 posterior on the 2026 snapshot pulls upstream toward the empirical rate (data dominates: 6{,}000 dam-months yields posterior 0.39% against prior 0.50%) and the other classes toward the literature prior (no observed positives, prior dominates).

The full level structure, the bound on each level’s contribution, and the data source feeding each level are summarised in **Table 5**. The design principle is that every level above the literature prior is *bounded* — no single operator, engineering feature, or InSAR reading can move a dam’s log-odds by more than a fixed amount — so that the single linked positive event cannot produce a spurious outlier.

4.3 Validation protocol

The pre-registered protocol in `docs/04-methods.md` calls for: (1) retrospective held-out positives at the two reference failures, with the model’s pre-failure probability trajectories as the headline qualitative result; (2) rolling-origin time-forward cross-validation with quarterly step; (3) leave-one-operator-out CV; (4) clustered bootstrap CIs over operators. The present manuscript reports training metrics for the population ranking, plus the two qualitative held-out retrospectives (§5.4, §5.5); full rolling-origin and operator-out CV await additional linked positives, which require the historical-SIGBM recovery described in §8.

4.4 Reproducibility

All code, data, scripts, processed outputs, and the visualisation are at github.com/ViktorSmirnov71/sentinela-bayes-br. The repository is designed for one-command reproduction of every artefact in this paper:

```
git clone https://github.com/ViktorSmirnov71/sentinela-bayes-br
cd sentinela-bayes-br
python -m venv .venv && source .venv/bin/activate
pip install -e ".[dev]"

python data/scripts/clean_sigbm.py
python data/scripts/build_cohort.py
python experiments/01_first_prediction/run.py
python data/scripts/export_viz_data.py
cd viz && python3 -m http.server 8765 # then open localhost:8765
```

5. Results

The cohort composition that the model operates over is summarised in **Figure 1** and **Table 1**. The 911-dam population is dominated by single-stage structures (53 %) and concentrated in Minas Gerais (35 %); upstream-method dams — the historically dangerous class — are only 45 facilities (5 %) but carry, in Table 1, the largest mean height (49 m vs. 12–22 m for other classes) and median impounded volume (4.4 Mm³ vs. ≤ 0.66 Mm³).

5.1 Construction-method posterior

Posterior 12-month failure rate per construction method (snapshot 2026-05-21) is given in **Table 2** and plotted in **Figure 2**:

construction	prior (annual)	empirical	posterior	n (dam-months)
upstream	0.50%	12 / 5,940 = 0.20%	0.39%	5,940
unknown	0.10%	0 / 1,056	0.090%	1,056
centerline	0.10%	0 / 15,180	0.040%	15,180
single_stage	0.10%	0 / 64,152	0.013%	64,152
downstream	0.05%	0 / 29,436	0.013%	29,436

This recapitulates the central published finding [arana2021] that upstream-method dams carry an order of magnitude higher historical failure rate than other construction methods, and quantifies the Brazilian-cohort- specific posterior given the single linked positive event. The upstream posterior (0.39 %/yr) is pulled below its prior (0.50 %) by the empirical rate of 0.20 %; the four other classes have no observed positives and sit near their priors.

5.2 Within-upstream ranking

Once the hierarchical Level-2 (operator) and Level-3 (height / volume / age / CRI) refinements are applied, the top-30 dams are no longer tied at the homogeneous upstream rate (**Figure 3**, **Table 3**). Selected entries from the top-10 (these are the pre-InSAR engineering-only values; dam_id 8765 falls to 1.60 % after its Level-4 InSAR shift, see §5.4):

rank	dam_id	name	operator	risk_12m
1	8765	Barragem de Germano	Samarco	1.97%
2	8332	Pontal	Vale	1.14%
3	9533	ED Monjolo	Vale	0.88%
4	9532	ED Vale das Cobras	Vale	0.63%
5	8431	Barragem Eustáquio	Kinross	0.63%
6	8286	Forquilha II	Vale	0.61%
7	8283	Forquilha I	Vale	0.59%
8	8290	Forquilha III	Vale	0.46%
9	8955	Barragem Serra Azul	ArcelorMittal	0.45%
10	8389	Sul Superior	Vale	0.42%

Independent validation against ANM emergency declarations. The model is given no information about ANM’s declared emergency levels. Yet the top-15 (Table 3) contains, purely from engineering features:

- **Barragem Serra Azul** (ArcelorMittal) at rank 9 — the *single dam in all of Brazil currently at Emergency Level 3*, the highest level.
- **Forquilha I, II and III** (Vale) at ranks 6–8 — all three at Emergency Level 2, the cluster that required partial evacuation of downstream communities in 2022.

- **Pontal, Sul Superior, Xingu** (Vale) — further Emergency Level 1–2 facilities.

That the model’s independent engineering-feature ranking surfaces the operationally-recognised emergency dams near the top is a meaningful external check: the ranking is not merely re-stating “upstream method” but is correctly ordering *within* the upstream class in a way that agrees with the regulator’s separate, monitoring-based determinations.

We caveat strongly that the *within-upstream* differentiation reflects the published engineering effect sizes (taller, larger, older dams fail more often) rather than any data-driven discovery; the InSAR Level-4 term refines this ranking with *current* deformation signal where Sentinel-1 coverage and our sampling sophistication permit (§5.4–5.5), but does not yet override the *historical* effect-size priors at population scale.

5.3 Calibration and discrimination

The model attains training log-loss 0.00064, AUROC 0.9995, ECE 0.00013. These figures are flattering because the cohort base rate is 0.01% per dam-month and the model assigns near-zero probability to the vast majority of rows. Substantive evaluation awaits more linked positives and the rolling-origin protocol.

5.4 Fundão retrospective: a partial-precursor finding

We submitted a second HyP3 batch (42 SBAS pairs) at the actual Fundão dam coordinates (-20.193, -43.493) over 2014-10 to 2015-11, downloaded the 42 unwrapped-phase rasters, sampled the line-of-sight time series at the dam centroid and at a 3 km stable reference point, and ran the hierarchical model on a rolling 12-month trailing-feature window for every month from 2014-10 to 2015-11.

Result, in full: ASF returned scenes only from June 2015 onward for this descending-orbit footprint (Sentinel-1’s systematic acquisition over MG began later in this orbit cycle than over neighbouring scenes), so the trailing-window InSAR features are active only from 2015-06 onward. For the months 2014-10 to 2015-05 the model returns the engineering-features baseline of **0.75 %**, correctly elevated relative to the 0.01 % cohort mean. From 2015-06 onward the predicted risk oscillates between **0.17 % and 3.27 %** across consecutive months. The month containing the catastrophic failure (2015-11) carries a predicted risk of **0.96 %**, slightly above the engineering baseline but not a clean accelerating-toward-failure trajectory ([figures/fundao_retrospective.png](#)).

The oscillation is informative about our InSAR sampling sophistication rather than about the underlying signal. Inspection of the raw time series shows that the dam-crest LOS values track the 3 km reference values to within ~1 mm — i.e. atmospheric and orbital common-mode signals dominate the differential, and the dam-specific motion in our crude median-pixel-buffer sampler is below the noise floor. The ~30 mm horizontal pre-failure deformation reported by @grebby2021 at Brumadinho (2019) was extracted via PS-InSAR with hundreds of reference points and explicit atmospheric correction; reproducing that precision is a methodology upgrade we have not yet built, and is the highest-priority follow-up.

What this experiment establishes:

1. **The pipeline works end-to-end on real ASF/HyP3 data.** Search, submit, queue, download, raster-sample, feature-extract, model-infer, plot — the chain has zero broken links.
2. **The model produces a plausible engineering-feature baseline** at the Fundão coordinates (0.75 %), well above the cohort mean, even when the dam is not in the SIGBM

registry.

3. **The current InSAR feature surface is too noisy** to materially refine the engineering-feature prediction at single-dam resolution. This is a known limitation of single-pixel-buffer sampling with no atmospheric correction and matches the methodological literature.

The negative-or-noisy retrospective is a research outcome, not a failure. It quantifies what’s missing from our pipeline (PSI processing, atmospheric correction, multi-reference de-trending) and pushes those items to the top of the methodology backlog. **Figure 5** (left) plots the trajectory with a 3-month smooth; **Figure 6** (left) plots the underlying LOS velocity. A side-by-side comparison of both retrospective events is given in **Table 4**.

5.5 Brumadinho B1 retrospective: a partial-precursor hit

A second HyP3 batch was submitted at the Brumadinho B1 dam coordinates (-20.117, -44.123) over 2018-01 to 2019-01. The 2018 era covers both Sentinel-1A and S1B operations, so the bbox + window yielded 66 scenes (vs. 16 for Fundão’s earlier-mission window), and 192 SBAS jobs were submitted at three-temporal-neighbour density.

Methodology finding during submission. 128 of 192 jobs failed at HyP3 with 100 % correspondence between failure and cross-relative-orbit pairing. Sentinel-1 has many tracks that overlap a small bbox, and pairs across tracks exceed the InSAR coherence baseline limit. We patched `short_baseline_pairs` to group by `pathNumber` before forming pairs; all future submissions group correctly. The 64 succeeded pairs are sufficient for a retrospective; the patch commit hash is [9fa0bbb](#).

Trajectory. With the rolling-window pipeline of §5.4 applied to the 64-pair B1 dataset, Sentinela’s predicted 12-month failure probability through 2018 evolves as follows (`figures/b1_brumadinho_retrospective.`

month	n_obs	risk_12m	notes
2018-01	3	0.42 %	engineering baseline; too-sparse window
2018-02	8	0.09 %	early-window noise
2018-03	13	0.09 %	early-window noise
2018-04	18	1.86 %	spike within Grebby’s first risk milestone
2018-05	23	0.57 %	
2018-06	28	0.35 %	drift downward
2018-07	33	0.13 %	
2018-08	39	0.22 %	
2018-09	43	0.26 %	
2018-10	49	0.28 %	
2018-11	54	0.23 %	
2018-12	59	0.20 %	
2019-01	61	0.26 %	collapse month

@grebby2021 identified two risk milestones for B1: a “definitive but emergent” window Feb–Aug

2018, and an “imminent risk of collapse” window Jun–Dec 2018. Our 2018-04 spike of 1.86 % falls squarely within the first window — a partial retrospective replication. The model does not, however, surface the second imminent-risk window; risk drops to baseline through summer/autumn 2018 and the collapse month carries only 0.26 %.

This is a meaningful partial result: the pipeline runs end-to-end on the canonical reference failure, produces a precursor flag in the expected calendar quarter, but does not sustain the signal through the second flagged window. As at Fundão, the limiting factor is the sampling sophistication of our LOS time series — single-pixel-median sampling against a 3 km reference cannot reproduce the PS-InSAR multi-reference processing Grebby used. Future iterations with MintPy will revisit both retrospectives. **Figure 5** (right) shows the B1 trajectory with the Grebby risk windows as amber bands; the smoothed April peak aligns with the leading edge of milestone 1.

5.6 Visualisation

A self-contained interactive Three.js visualisation is shipped at `viz/index.html`: 877 geocoded dams are rendered as risk-proportional spikes erupting from a 3D mesh of Brazil’s terrain (SRTM elevation via Copernicus terrarium tiles), with colour along a teal → violet → magenta ramp by risk decile, the top-30 pulsing, and active-emergency dams carrying an amber ground halo. **Figure 4** is a static render of the same risk field for print.

The 3D view surfaces patterns the tabular ranking does not: the dense clustering of high-risk spikes in the iron-ore corridor of the Minas Gerais highlands (terrain elevation ~800–1,400 m), the relative emptiness of the Amazon basin and the northeast, and the spatial proximity of the Forquilha cluster to dam_id 8765 at the Mariana / Germano complex. Future iterations will encode posterior uncertainty as faded outer spikes.

6. Discussion

6.1 What we learn

The single most useful finding of the present iteration is methodological: *with one linked positive event in a 911-dam cohort, the literature priors are the model*. No amount of feature engineering or model-class tuning can extract more signal than the data contains. The honest path is to embed the literature transparently and refine the priors as linked positives accumulate.

Specific within-Brazil findings, with all the caveats above:

1. **Upstream-method posterior is consistent with literature.** Our Level-1 posterior at 0.39%/year for upstream-method dams is within the range reported in @rana2021’s global analysis (0.4–1.2%/year). The construction-method effect is the single dominant predictor.
2. **Vale dominates the top-of-shortlist.** Six of the top-10 dams are Vale-operated. This reflects Vale’s large fleet of upstream-method iron-ore dams in Minas Gerais; the model does *not* infer an operator-level penalty for Vale beyond its dam count.
3. **The Forquilha cluster’s appearance in the top-10 validates the engineering-feature surface.** ANM’s 2022 Emergency Level 2 declaration on these three dams was based on geotechnical monitoring outside our feature set; our model’s independent agreement is evidence the model’s ranking is operationally interpretable.

6.2 What we do not learn

- **Whether InSAR adds marginal predictive value.** This is RQ1 in our pre-registration. Pending HyP3 completion (currently 42 jobs running).
- **Whether B1 / Brumadinho would have been flagged prospectively.** This is RQ4. Requires a pre-2019 SIGBM snapshot (Wayback or FOI).
- **Whether ANM's CRI score adds information beyond engineering features.** RQ5. Requires more linked positives for meaningful comparison.

6.3 Ethics

This work is *research*, not an operational early-warning system. Outputs must not be used to direct evacuation decisions. The model uses only public data; operator-internal monitoring signals (piezometers, inclinometers) are deliberately excluded, both to preserve open reproducibility and to lower-bound what the public-data approach can achieve. We will not publish a public dashboard naming individual operators without giving them an opportunity to respond. Full ethics posture: docs/06-ethics-and-limitations.md.

7. Limitations and threats to validity

1. **Registry amnesia.** Failed dams are removed from the active SIGBM post-failure. This is the dominant bias in the training data.
2. **One linked positive.** Within-cohort signal at the engineering-feature level is dominated by the literature priors, not by the data.
3. **Survivor selection.** Dams that approached failure but were remediated are coded as non-failures in our labels.
4. **Sample selection on satellite era.** All InSAR-era failures (2014–) may not generalise to pre-radar regimes; we do not claim such generalisation.
5. **Label provenance.** Bowker–Chambers severity coding is one analyst's judgement. Sensitivity analyses at severity- ≥ 3 are planned.

8. Future work

In priority order, informed by the results above:

1. **PS-InSAR / MintPy processing.** The single highest-value upgrade. Both retrospectives are limited by atmospheric noise in single-pixel sampling; multi-reference persistent-scatterer processing should lift the second Grebby window above the noise floor and is the prerequisite for InSAR to materially improve the population ranking.
2. **Historical SIGBM snapshots.** Wayback Machine archives plus an FOI request to ANM should yield pre-2015 and pre-2019 cohort states so that Fundão and B1 can be predicted prospectively, not via proxy, and so the model gains genuinely held-out positives.
3. **CHIRPS rainfall.** Daily 0.05° rainfall + SPI-3 / SPI-6 antecedent moisture as the next feature block; enables the rainfall $\times$ InSAR interaction test (RQ7).
4. **Operator-level Bayesian hierarchy.** The current James–Stein shift is a first-order operator effect; a full random-effects model with PyMC will be fit once linked positives ≥ 30 .
5. **Cross-LATAM transfer.** SIGBM's analogues exist for Colombia (DANE) and Mexico (INEGI). Cross-country transfer experiments will test the generality of the model's structure beyond Brazil.

9. Conclusion

We have built and openly released Sentinela, a calibrated hierarchical-Bayesian forecaster for Brazilian mine-tailings dam failures driven entirely by public data. The system is reproducible from a clean clone to forecast in five commands, and every artefact in this paper — the 911-dam cohort, the per-dam risk ranking, the 3D risk field, and the two retrospectives — regenerates deterministically.

Three findings stand out. First, in the one-linked-positive regime that characterises this problem, discriminative learners fail (they memorise or smooth to zero) and a literature-prior Bayesian model is the honest choice; its construction-method posterior reproduces the published order-of-magnitude upstream-method risk premium. Second, the model's purely engineering-feature ranking independently surfaces the dams ANM has separately placed under emergency declarations — the single Level-3 dam in the country and the entire Level-2 Forquilha cluster appear in the top-15 — which is a genuine external validation that the within-class ordering is meaningful. Third, on the canonical Brumadinho B1 failure, the model produces a risk spike in the exact calendar quarter of the precursor window reported independently by satellite geodesists, though it does not sustain that signal — a partial replication that cleanly localises the remaining methodological gap to InSAR processing sophistication rather than to the modelling framework.

Sentinela is a research artefact, not an operational early-warning system, and its limitations (registry amnesia, one linked positive, single-pixel InSAR sampling) are documented and bounded rather than hidden. What it demonstrates is that a small, carefully-chosen open-data pipeline plus a transparent Bayesian model can produce a calibrated, externally-validated, fully-reproducible risk ranking for a high-stakes infrastructure problem — and can do so quickly enough to be a credible decision-support input rather than a multi-year research programme.

Figures and tables

Figures (in figures/):

- Figure 1 — `fig1_cohort_composition.png`: cohort by method / state / emergency / status.
- Figure 2 — `fig2_posterior_rates.png`: prior vs. posterior failure rate per method.
- Figure 3 — `fig3_top_risk_ranking.png`: top-20 dams by predicted 12-month risk.
- Figure 4 — `fig4_terrain_3d.png`: static 3D render of the risk field over terrain.
- Figure 5 — `fig5_retrospectives.png`: Fundão + B1 risk trajectories with Grebby windows.
- Figure 6 — `fig6_insar_comparison.png`: extracted InSAR LOS velocity per event.

Tables (in paper/tables/, CSV + tables.md):

- Table 1 — cohort summary by construction method.
- Table 2 — prior vs. posterior failure rate.
- Table 3 — top-15 highest-risk dams.
- Table 4 — retrospective comparison (Fundão vs. B1).
- Table 5 — hierarchical model structure.

Acknowledgements

This is pro bono research conducted on behalf of [Arakon](#) as a theoretical thought experiment on rapid machine-learning model deployment for high-stakes infrastructure decisions. The author thanks ANM for maintaining SIGBM as public open data, ESA/Copernicus for Sentinel-1 open archive, and Alaska Satellite Facility for hosting HyP3.

References

References use BibTeX keys defined in `docs/refs.bib`. Pandoc invocation:

```
pandoc paper/manuscript.md \
  --citeproc \
  --bibliography docs/refs.bib \
  --csl=https://www.zotero.org/styles/nature \
  -o paper/manuscript.pdf
```

Primary references cited above:

- Grebby et al. 2021 — *Nat. Comms Earth & Env.* — Brumadinho precursors.
- Carlà et al. 2019 — *Sci. Rep.* — InSAR inverse-velocity.
- Du et al. 2020 — *Sci. Total Env.* — Brumadinho InSAR risk.
- Gama et al. 2023 — *Nat. Comms Earth & Env.* — slip-surface mechanism.
- Mura et al. 2018 — *Remote Sens.* — Germano post-Fundão DInSAR.
- Macchiarulo et al. 2023 — *Remote Sens. Env.* — joint InSAR + moisture.
- Manconi et al. 2024 — *Bull. Eng. Geol.* — operational Sentinel-1.
- Zhang et al. 2025 — *Int. J. Appl. Earth Obs.* — consolidation separation.
- Zheng et al. 2025 — *Sci. Rep.* — InSAR + numerical model.
- Rana et al. 2021 — *Eng. Geol.* — global TSF failure statistics.
- de Lima et al. 2022 — SIGBM risk classification.
- ANM SIGBM (app.anm.gov.br/sigbm/publico).
- Bowker & Chambers, WMTF database.
- Ferreira et al. 2020 — BrazilDAM benchmark.

Appendix A — Figures

Figure 1 — Brazilian mine-tailings cohort composition (SIGBM 2026, n=911)





Figure 2 — Prior vs. posterior failure rate by construction method






Figure 4 — Predicted failure-risk field over Brazilian terrain (spike height scales with 12-month failure probability; only top-risk dams drawn)




Figure 5 — Retrospective risk trajectories at the two reference failures (amber bands = Grebby 2021 risk windows)





Figure 6 — Extracted InSAR line-of-sight velocity across the rolling window



Appendix B — Tables

Table 1 — Cohort summary by construction method

construction_method	n_dams	mean_height_m	median_volume_m3	n_high_cri	n_emergency
single_stage	486	11.9	7.657e+04	29	38
downstream	223	22.1	6.632e+05	14	14
centerline	115	22.4	3.197e+05	17	19
upstream	45	49.1	4.42e+06	13	14
unknown	42	2.2	0	4	4

Table 2 — Prior vs. posterior failure rate (Beta-Binomial, $\alpha=10,000$)

construction_method	n_dam_months	positives	prior_annual_pct	posterior_annual_pct
upstream	5940	12	0.5	0.389
unknown	1056	0	0.1	0.09
centerline	15180	0	0.1	0.04
single_stage	64152	0	0.1	0.013
downstream	29436	0	0.05	0.013

Table 3 — Top-15 highest-risk active dams (2026 snapshot)

dam_id	name	operator_name	state	construction_method	height_m	emergency_risk	risk_12m_pct
8765	Barragem de Germano	SAMARCO MINERA-CAO S.A.	MG	upstream	163	0	1.595
8332	Pontal	VALE S.A.	MG	upstream	69.9	1	1.143
9533	ED Monjolo	VALE S.A.	MG	upstream	145	0	0.879

dam_id	name	operator_name	state	construction	height_m	emergency_risk	risk_12m_pct
9532	ED Vale das Cobras	VALE S.A.	MG	upstream	122.4	0	0.635
8431	Barragem Eustáquio	KINROSS BRASIL MINERACAO S/A	MG	centerline	114	0	0.632
8286	Forquilha II	VALE S.A.	MG	upstream	94.3	2	0.615
8283	Forquilha I	VALE S.A.	MG	upstream	94.44	2	0.587
8290	Forquilha III	VALE S.A.	MG	upstream	77	2	0.464
8955	Barragem Serra Azul	ARCELORMITTAL BRASIL S.A.	MG	upstream	80.1	3	0.453
8389	Sul Superior	VALE S.A.	MG	upstream	75.5	2	0.422
9534	Xingu	VALE S.A.	MG	upstream	72	1	0.406
8209	Campo Grande	VALE S.A.	MG	upstream	93.98	0	0.406
9320	Barragem B2 - Mina Tico-Tico	MINERACAO MORRO DO IPE S.A.	MG	upstream	86	0	0.399
8291	Barragem MSG	MINERACAO SERRA GRANDE S A	GO	upstream	92	0	0.388
8207	Doutor	VALE S.A.	MG	upstream	79.98	0	0.342

Table 4 — Retrospective comparison at the two reference failures

event	collapse_date	insar_pairs	max_risk_pct	collapse_month	peak_risk_month	matches_grebbby_window
Fundão (Samarco)	2015-11-05	42	3.27	0.96	2015-06	n/a (Fundão not in Grebby)
Brumadinho B1 (Vale)	2019-01-25	61	1.86	0.26	2018-04	yes (peak in milestone-1)

Table 5 — Hierarchical model structure

level	mechanism	bound	data_source
0 prior	literature base rate (Rana 2021 / Bowker-Chambers)	—	literature
1 construction	Beta-Binomial posterior per construction method	—	SIGBM + WMTF labels
2 operator	James-Stein shrunk per-operator logit shift	± 1.0 logit	SIGBM + labels
3 engineering	bounded logit shifts: height, volume, age, CRI	0.10–0.30 per z	SIGBM static fields
4 InSAR	bounded logit shifts: LOS velocity, accel, spectral slope, variance ratio	± 1.5 logit total	Sentinel-1 HyP3 InSAR